EVAN LI

Designing, shipping, and breaking agentic systems

Boston, MA · 914.486.8529 · evanlie.737@gmail.com

evanli.org · linkedin.com/in/evan-li-ff · github.com/evanli135 · evanli.org/scholar

PUBLICATIONS

ACE: A Security Architecture for LLM-Integrated App Systems NDSS 2026

Evan Li, Tushin Mallick, Evan Rose, William Robertson, Alina Oprea, Cristina Nita-Rotaru *Co-First Author*

Top-4 security venue · 77+ citations · arxiv.org/abs/2504.20984

EXPERIENCE

Software Engineer, R&D Networks (Internship) Jun 2026 – Aug 2026

Keysight Technologies *Calabasas, CA*

- Built an end-to-end agentic penetration testing service that automates SQL injection, prompt injection, and DoS attacks against firewall devices
- Engineered an MCP service giving LLM agents guarded control of physical test chassis, with persistent user preferences and RAG over the BreakingPoint attack corpus; outperformed the prior MCP baseline by 41% on utility benchmarks
- Prototyped an AI anomaly classification pipeline for a network security and stress emulation platform across L4–L7 topologies, automating health checks and max-throughput/breaking-point detection

Software Engineer, Agentic AI (Co-op) Jan 2026 – Jun 2026

Vestmark *Cambridge, MA*

- Built safeguards against unauthorized actions by financial advisory agents including human-in-the-loop approval gates for portfolio modifications and Amazon Bedrock Guardrails blocking leakage of sensitive company data
- Designed memory storage, context retrieval, persistence, and observability (OpenTelemetry/Grafana) layers for agentic news ingestion pipelines in TypeScript and Mastra
- Parallelized agentic tooling surface, improving API latency by 67%

Agentic AI Security Researcher (Co-op) Jan 2025 – Jun 2025

Khoury Network and Distributed Systems Security Lab *Boston, MA*

- Co-designed ACE, an agent security architecture that plans only from trusted information and enforces safe execution via static analysis and rule-based checks
- Built ACE's core components (context management, abstract planner, static checker, execution sandboxes) and evaluation pipelines using Python, Docker, and LangChain
- Evaluated ACE on InjecAgent, Agent Security Bench, and newly discovered planning and execution attacks; blocked 100% of evaluated prompt injection attacks and verified system-level safety guarantees
- Red-teamed IsolateGPT (NDSS 2025), uncovering critical prompt injection and DoS vulnerabilities that invalidate its claimed app isolation guarantees

NLP Engineering Intern (Internship) Jul 2024 – Sep 2024

Matrix Origin *San Francisco, CA*

- Built evaluation infrastructure for an NL-to-SQL system; benchmarked Llama across hundreds of edge-case prompts
- Improved query accuracy 40% via systematic prompt engineering and few-shot example construction

PROJECTS

Andal: Embedded Event Store C / Python

- Built a database engine from scratch in C with columnar storage, write-ahead log, crash recovery, and inverted indices
- Designed for embedded analytics: immutable append-only segments, lazy loading, and a Python C API: pip install andal

EDUCATION

B.S. Computer Science (Systems Engineering Concentration) Sep 2023 – May 2027

Northeastern University, GPA: 3.81 *Boston, MA*

Coursework: Distributed Systems, Computer Systems, Cybersecurity, Networks, Databases, Algorithms, Machine Learning

SKILLS

AI Security: Prompt Injection Defense, Multi-Agent Threat Modeling, Context Sandboxing, Information Flow Control
Languages: Python, C, Rust, C++, TypeScript, JavaScript, SQL, Bash, Ruby, Java
AI/ML: LangChain, Mastra, MCP, Agent Orchestration, LLM Evaluation, Prompt Engineering, Scikit-Learn
Infrastructure: Docker, Kubernetes, Linux, AWS, PostgreSQL, MongoDB, Redis, OpenTelemetry, CI/CD
Frameworks/Libraries: React, FastAPI, Ruby on Rails, Django, Express, Pandas, NumPy, TailwindCSS